

# Module 4 - Process scheduling

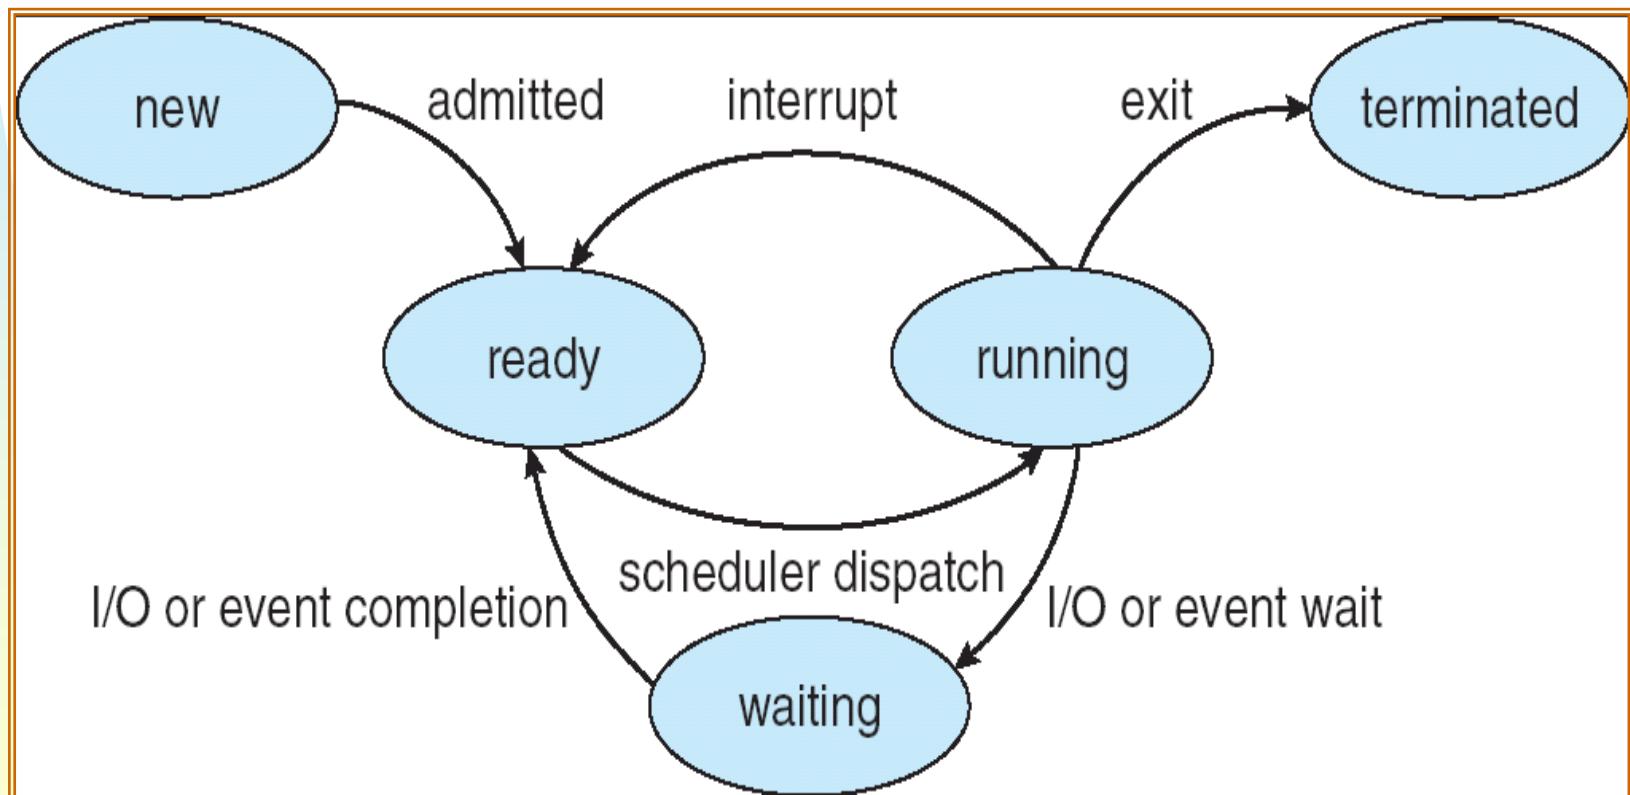
---

**Reading: Chapter 6 (Silberschatz)**

# Module overview

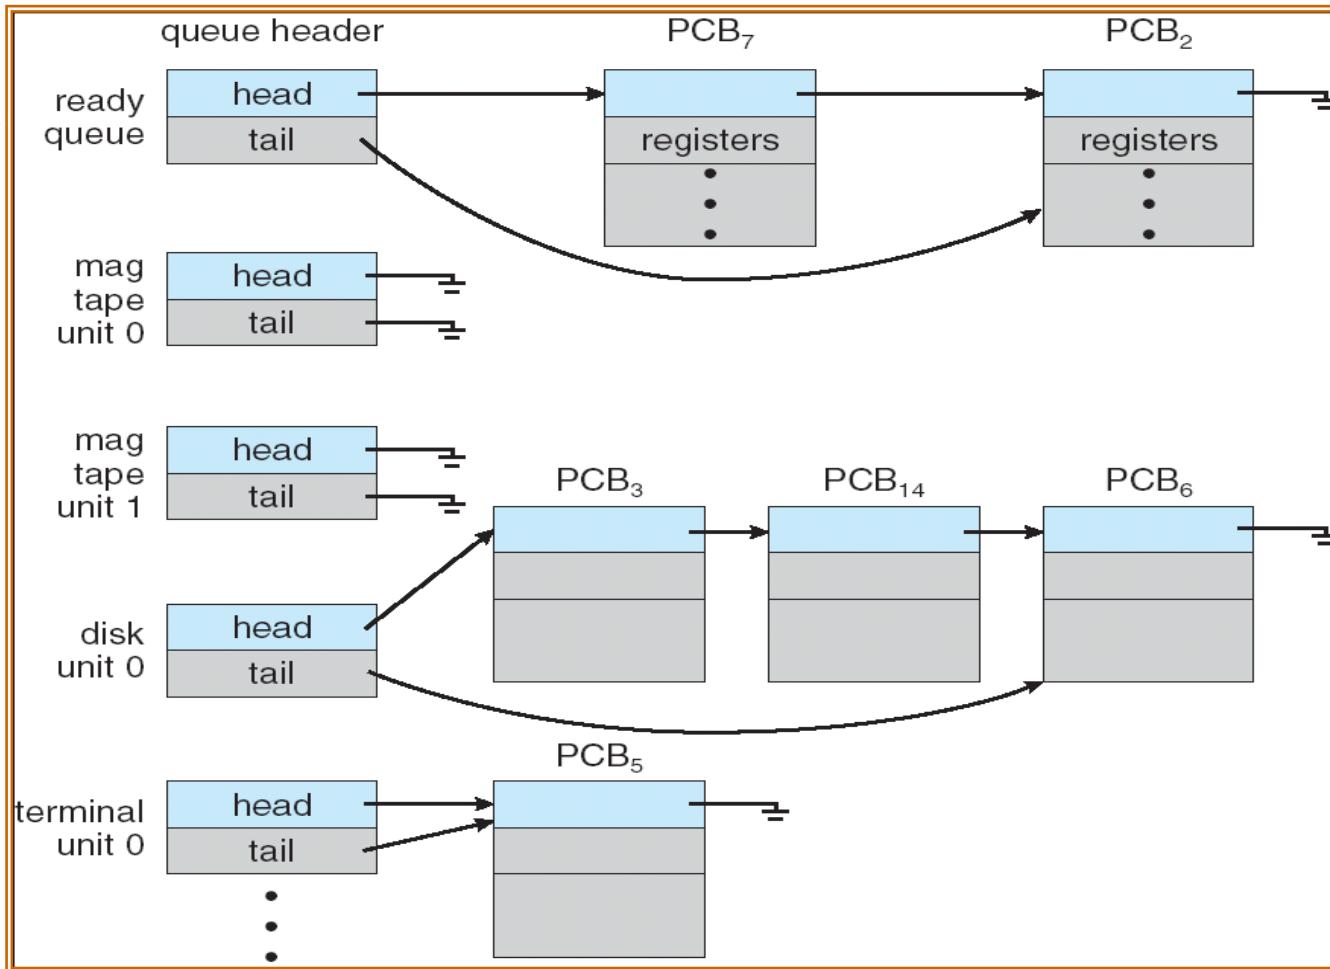
- **Basic concepts**
- **Scheduling criteria**
- **Scheduling algorithms**
- **Multiprocessor scheduling**
- **Algorithm evaluation**

# State transition diagram of a process



# Process queues for scheduling

Ready queue

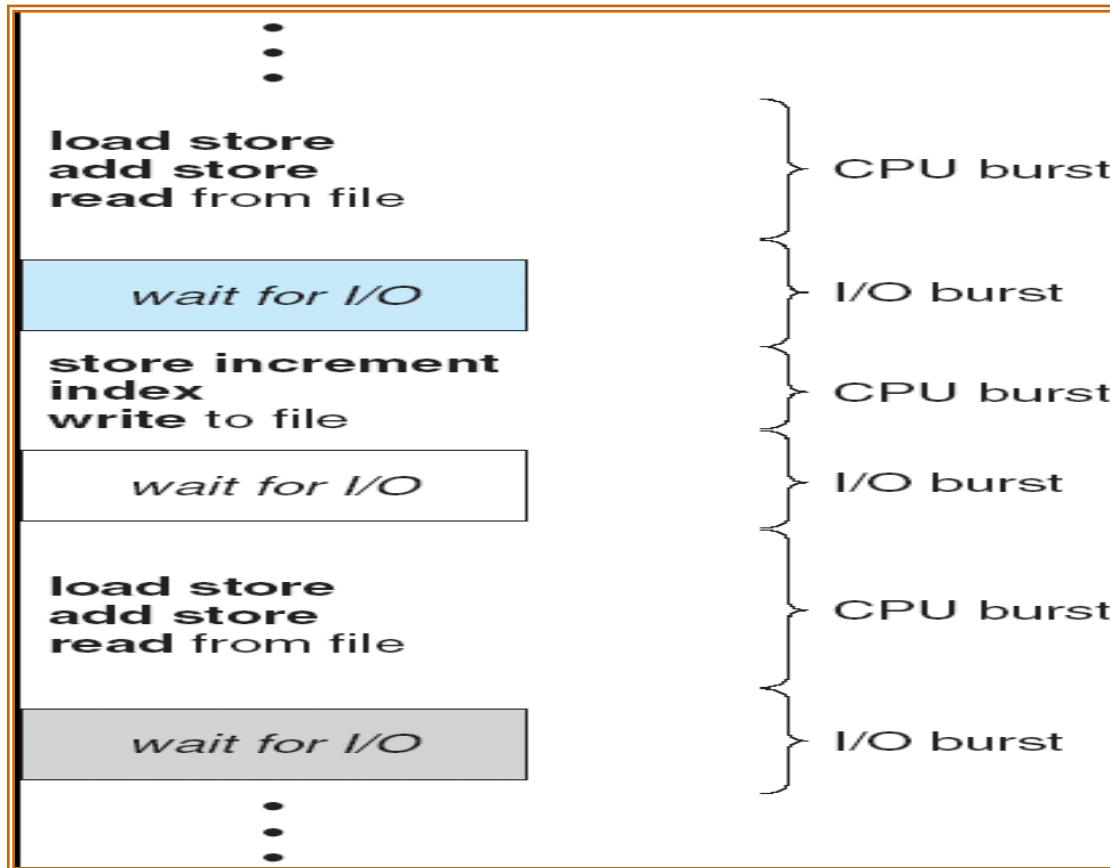


We will assume that the first process in a queue is the one that uses the resource: here, proc7 executes

# Basic concepts

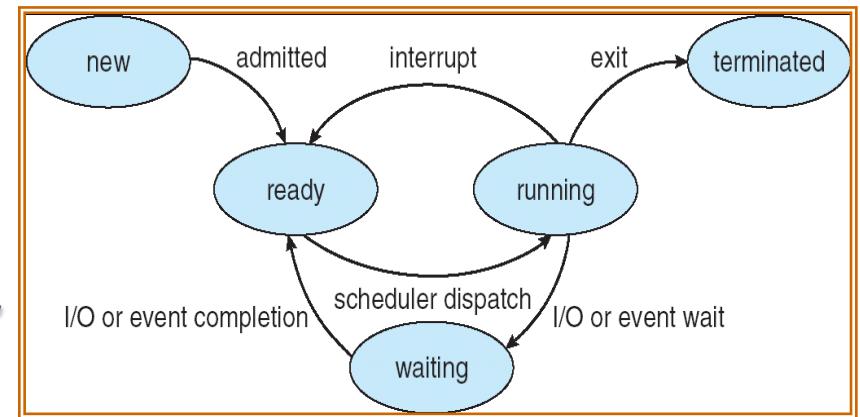
- **Multiprogramming is designed to achieve maximum use of resources, especially the CPU**
- **The CPU scheduler is the part of the OS that decides which process in the ready queue gets the CPU when it becomes free**
  - should aim for optimal use of the CPU
- **CPU is the most valuable resource in a computer, so we are talking about it**
  - However, the principles that we will see also apply to the scheduling of other resources (I / O units, etc.).
- **Must understand process behavior**
  - To make the right scheduling decision

# The cycles of a process



- **CPU and I / O cycles (bursts): the execution of a process consists of execution sequences on CPUs and I / O waits**

## When to invoke the scheduler



- **The scheduler must make its decision every time the executing process is interrupted, i.e.**
  - a process presents itself as **new** or **finished**
  - an executing process becomes **stuck waiting**
  - a process changes from **running** to **ready**
  - a process changes from **waiting** to **ready**
  - in conclusion, any event in a system causes an interruption of the CPU and the intervention of the scheduler, who will have to make a decision about which proc. or thread will have the CPU after
- **Preemption:** we have preemption in the last two cases if we remove the CPU from a process that had it and can continue to use it
- In the first two cases, there is **no preemption**
- Several problems to resolve in the case of preemption

# Dispatcher

- The code of the OS that gives control to the process chosen by the scheduler. It must be concerned with:
  - change context
  - change to user mode
  - restart the chosen process
- Dispatcher latency
  - the time required to perform the duties of the dispatcher
  - it is often overlooked, it must be assumed that it is small compared to the length of a cycle

# Scheduling criteria

- There will normally be several processes in the ready queue
- When CPU becomes available, which one to choose?
- The general idea is to make the choice in the interest of the efficient use of the machine
- But the latter can be judged according to different criteria ...

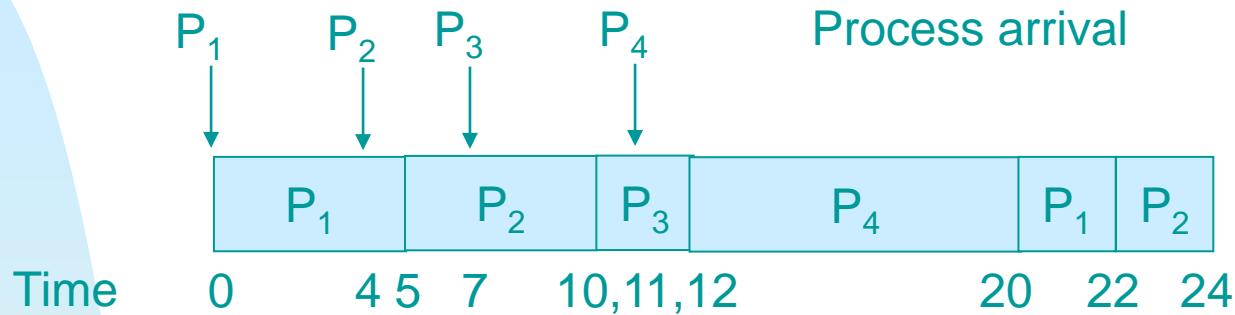
# Scheduling criteria

- **Main reason for scheduling**
  - Percentage of use: keep CPUs and I / O modules busy
- **Time-sharing systems?**
  - Response time (for interactive systems): the time between a request and the response
- **Servers?**
  - Throughput: number of processes that complete in the unit of time
- **Batch processing systems?**
  - Turnaround time: the time taken by the process from its arrival to its end.
- **Loaded systems?**
  - Waiting time: waiting in ready queue (sum of all time spent in ready queue)

# Scheduling criteria: maximize / minimize

- **To maximize**
  - CPU use: percentage of use
  - Throughput: number of processes that complete in the unit of time
- **To minimize**
  - Response time (for interactive systems): the time between a request and the response
  - Turnaround time: completion time minus arrival time
  - Waiting time: waiting in ready queue

# Scheduling Criteria Example



- **CPU utilization:**
  - 100%
- **Throughput :**
  - 4/24
- **Turnaround time ( $P_3$ ,  $P_2$ ):**
  - $P_3$ : 5
  - $P_2$ : 20
- **Waiting time ( $P_2$ ):**
  - $P_2$ : 13
- **Response time ( $P_3$ ,  $P_2$ ):**
  - $P_3$ : 3
  - $P_2$ : 1

**Now let's take a look at several scheduling methods and see how they behave against these criteria.**

***we will study specific cases***

*the study of the general case would require recourse to probabilistic or simulation techniques*

# First come, first served (First come, first serve, FCFS)

- Note, no preemption

<u>Process</u>	<u>Cycle time</u>
P1	24
P2	3
P3	3

If the processes arrive at time 0 in the order: P1, P2, P3

The Gantt chart is:



Waiting time for P1 = 0; P2 = 24; P3 = 27

Average waiting time:  $(0 + 24 + 27) / 3 = 17$

# First come, first served

- CPU utilisation = 100%
- Throughput =  $3/30 = 0.1$ 
  - 3 processes completed in 30 units of time
- Average rotation time:  $(24 + 27 + 30) / 3 = 27$



# FCFS scheduling (continued)

If the same processes arrive at 0 but in order

$P_2, P_3, P_1$ .

The Gantt chart is:



- Waiting time for  $P_1 = 6$   $P_2 = 0$   $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3) / 3 = 3$
- Much better!
- So for this technique, the average wait time can vary greatly
- *Exercise: also calculate the average rotation time, throughput, etc.*

# Take into account the arrival time!

- If the processes arrive at different times, subtract the arrival times
- Exercise: repeat the calculations if:
  - P2 arrives at time 0
  - P1 arrives at time 2
  - P3 arrives at time 5

# Convoy Effect with the FCFS

- Consider a single CPU bound process and many I/O bound processes (fairly normal situation).
- The I/O bound processes wait for the CPU: I/O under utilized (\*).
- The CPU bound process requests I/O: the other processes rapidly complete CPU bursts and go back to I/O: CPU under utilized.
- CPU bound process finishes I/O, and so do the other processes: back to \*
- Solutions?

# FCFS Scheduling Discussion

**Is it simple and easy to program?**

- Yes!

**Can we find other benefits?**

- It costs little time to make scheduling decision

**Does it provide low waiting time?**

- Not at all, can be quite awful

**Does it provide good CPU utilization?**

- No – the convoy effect

**OK, forget about it.**

**Now, let's try to reduce the average waiting time**

**Idea: It is better to execute ready processes with short CPU bursts**

## Shorter First = Shortest Job First (SJF)

- **The shortest process starts first**
- **Optimal in principle in terms of average waiting time**
  - (see the last example)
- **But how do we know**

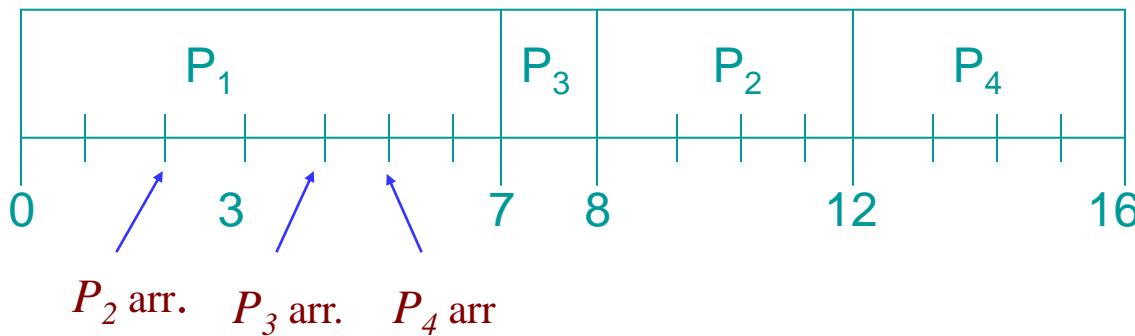
# Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
  - nonpreemptive – once CPU given to the process it cannot be preempted until completes its CPU burst
  - preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is also known as the Shortest-Remaining-Time-First (SRTF)
    - Just call it preemptive SJF
- Preemptive SJF is optimal – gives minimum average waiting time for a given set of processes
  - Moving a process with short CPU burst in front of a process with longer CPU burst reduces average waiting time

# Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (non-preemptive)

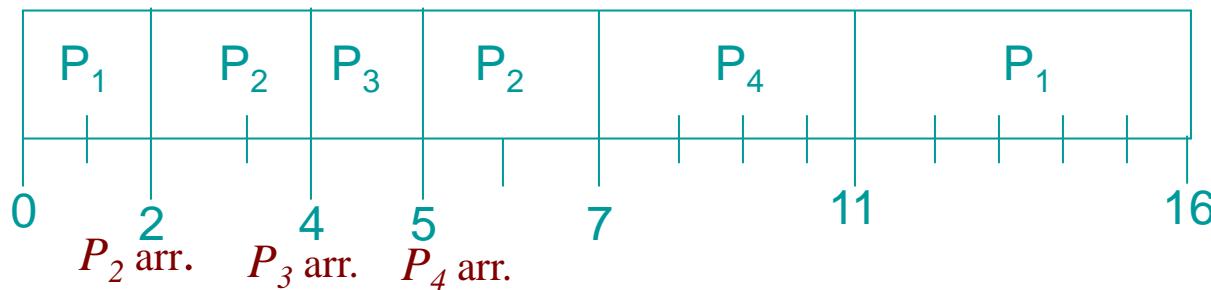


- Average waiting time =  $(0 + 6 + 3 + 7)/4 = 4$

# Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (preemptive)



- Average waiting time =  $(9 + 1 + 0 + 2)/4 = 3$

# Small technical detail with SJF

How do we know the length of the next CPU burst?

- If you know that, you probably know also tomorrow's stock market prices...
  - You will be rich and won't need to waste time in CSI3131 class
- So, we can only estimate it

**Any idea how to estimate the length of the next CPU burst?**

- Probably similar as the previous bursts from this process
- Makes sense to give more weight to the more recent bursts, not just straightforward averaging
  - Use exponential averaging

# Exponential averaging with SJF

1.  $t_n$  = actual length of  $n^{th}$  CPU burst
2.  $\tau_{n+1}$  = predicted value for the next CPU burst
3.  $\alpha, 0 \leq \alpha \leq 1$  = relative weight of recent vs past history
4. Define:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

$$\begin{aligned}\tau_{n+1} &= \alpha t_n + (1 - \alpha) \alpha t_{n-1} + (1 - \alpha)^2 \alpha t_{n-2} + \dots \\ &\quad + (1 - \alpha)^i \alpha t_{n-i} + \dots + (1 - \alpha)^n t_1\end{aligned}$$

# The shortest first SJF: review

- **Difficulty estimating the length in advance**
- **Long processes will suffer from *famine* when there is a constant supply of short processes**
- **Preemption is required for timeshared environments**
  - A long process can monopolize the CPU if it is the first to enter the system and it does not do I / O
- **There is an implicit assignment of priorities: preferences for shorter jobs**

# SJF Discussion

## Does it ensure low average waiting time?

- Yes, it was designed that way
  - As long as our burst-length predictions more-or-less work

## Does it provide low response time?

- Not necessarily, if there is a steady stream of short CPU bursts, the longer bursts will not be scheduled
- This is called starvation
  - A process is blocked forever, always overtaken by other processes (well, or at least while the system is busy)

Let's see Priority Scheduling.

# Priority Scheduling

- A priority number (usually integer) is associated with each process
  - On some systems (Windows), the higher number has higher priority
  - On others (Unix) , the smaller number has higher priority
- The CPU is allocated to the process with the highest priority
  - Can be preemptive or non-preemptive
  - but usually you want to preempt low-priority process when a high priority process becomes ready
- Priority can be explicit
  - Assigned by the sysadmin/programmer, often for political reasons
    - Professor jobs are of higher priority
  - But also for technical reasons
    - This device has to be serviced really fast, otherwise the vital data will be lost (real-time processing)
- Priority can also be implicit (computed by the OS)
  - SJF can be seen as priority scheduling where priority is the predicted next CPU burst time

# Priority Scheduling Discussion

## Good properties

- Professor jobs will be scheduled before student jobs
- Allows support of real-time processing

## Bad properties

- Professor jobs will be scheduled before student jobs
  - OK, give me something else
- starvation – low priority processes may never execute

## How to resolve the starvation problem?

- aging – keep increasing the priority of a process that has not been scheduled for a long time

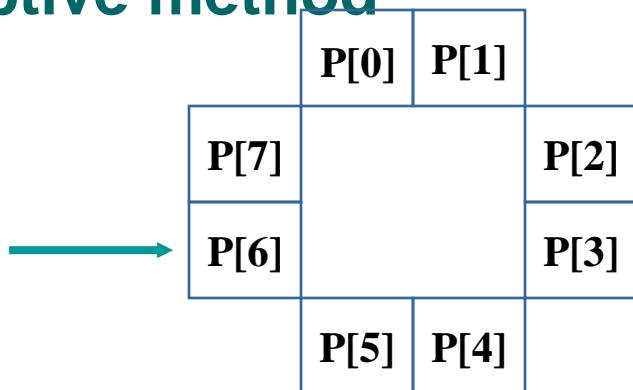
## What to do with the processes of the same priority level?

- FCFS
- Might as well add preemption = Round Robin

# Round-Robin (RR)

Most used in practice

- Each process is allocated a quantum of time (e.g. 10-100 millisecs.) To run
  - (book terminology: *time slice*)
- If it runs for an integer quantum without other interruptions, it is interrupted by the timer and the CPU is given to another process
- The interrupted process becomes ready again (at the end of the queue)
- Preemptive method



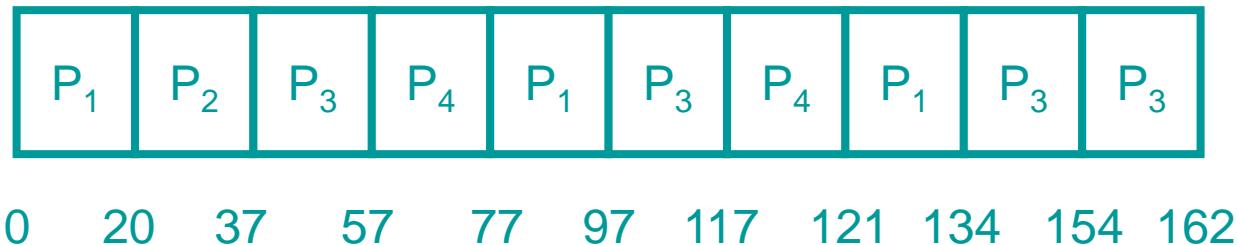
The ready queue is a circle (including RR)

# Round-Robin performance

- If there is a process in the ready queue and the quantum is  $q$ , then each process receives  $1 / n$  of the CPU time in units of max duration.  $q$
- If  $q$  large  $\Rightarrow$  FCFS
- If  $q$  small ... we will see

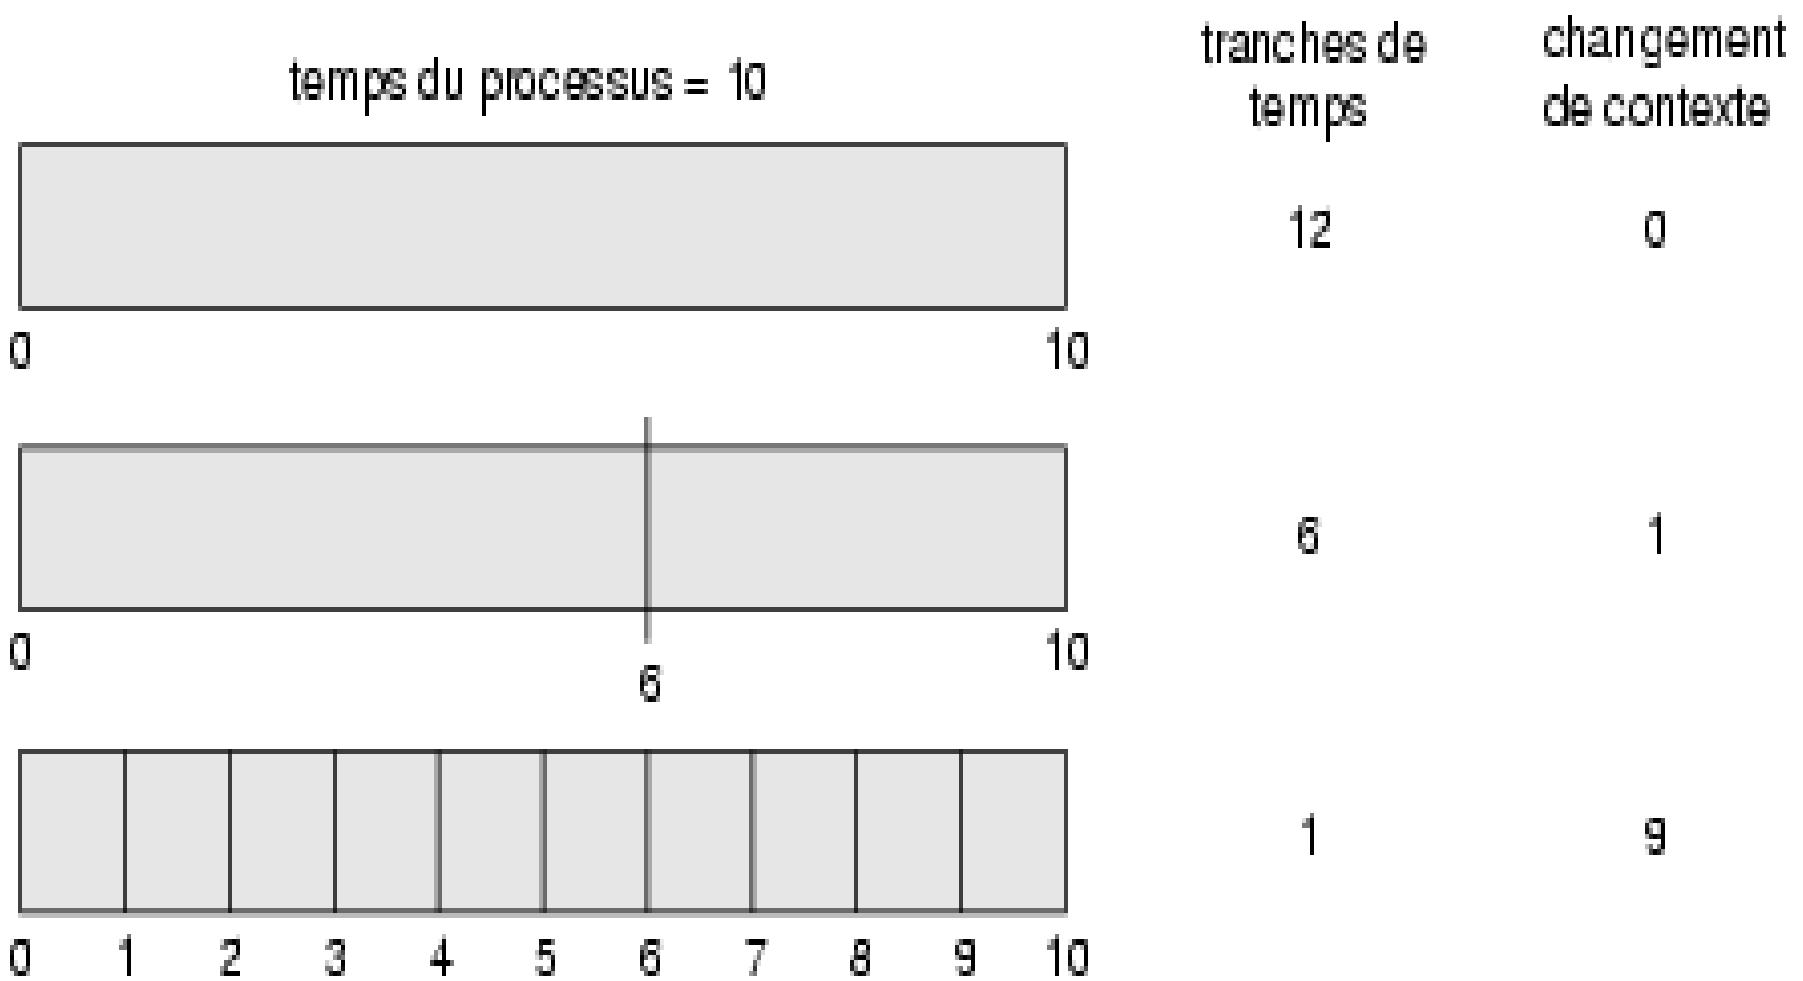
# Example: RR Quantum = 20

<u>Process</u>	<u>Cycle</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24



- **Normally,**
  - higher turnaround time than SJF
  - but average waiting time better

# A small quantum increases context switches (os time)



## Example to see the importance of a good choice of quantum (to be developed as an exercise)

- **Three cycles:**
  - A, B, C, all of 10
- **Try with:**
  - $q = 1$
  - $q = 10$
- **In this second case, the round-robin works as FCFS and the average rotation time is better**

# Algorithms we have seen so far

- **First Come First Serve**
  - simple, little overhead, but poor properties
- **Shortest Job First**
  - needs to know CPU burst times
  - exponential averaging of the past
- **Priority Scheduling**
  - This is actually a class of algorithms
- **Round Robin**
  - FCFS with preemption

# Scheduling exercises

Consider three processes P1, P2, P3

- Burst times for P1: 14,12,17
- Burst times for P2: 2,2,2,3,2,2,2,3,2,2,2,3,2,2,2,3
- Burst times for P3: 6,3,8,2,1,3,4,1,2,9,7
- All three arrive at time 0, in order P1, P2, P3
- Each CPU burst is followed by an I/O operation taking 6 time units
- Let's simulate the scheduling algorithms
  - FCFS
  - Round Robin (quantum=5)
  - Non-preemptive SJF or Preemptive SJF (your choice)
  - Round robin (quantum=5) with Priority scheduling, priorities are P2=P3>P1

# Multilevel Queue

Idea: **Partition the ready queue into several queues, and handle each queue separately.**

**Example:**

- foreground queue (interactive processes)
- background (batch processes)

**Each queue might have its own scheduling algorithm**

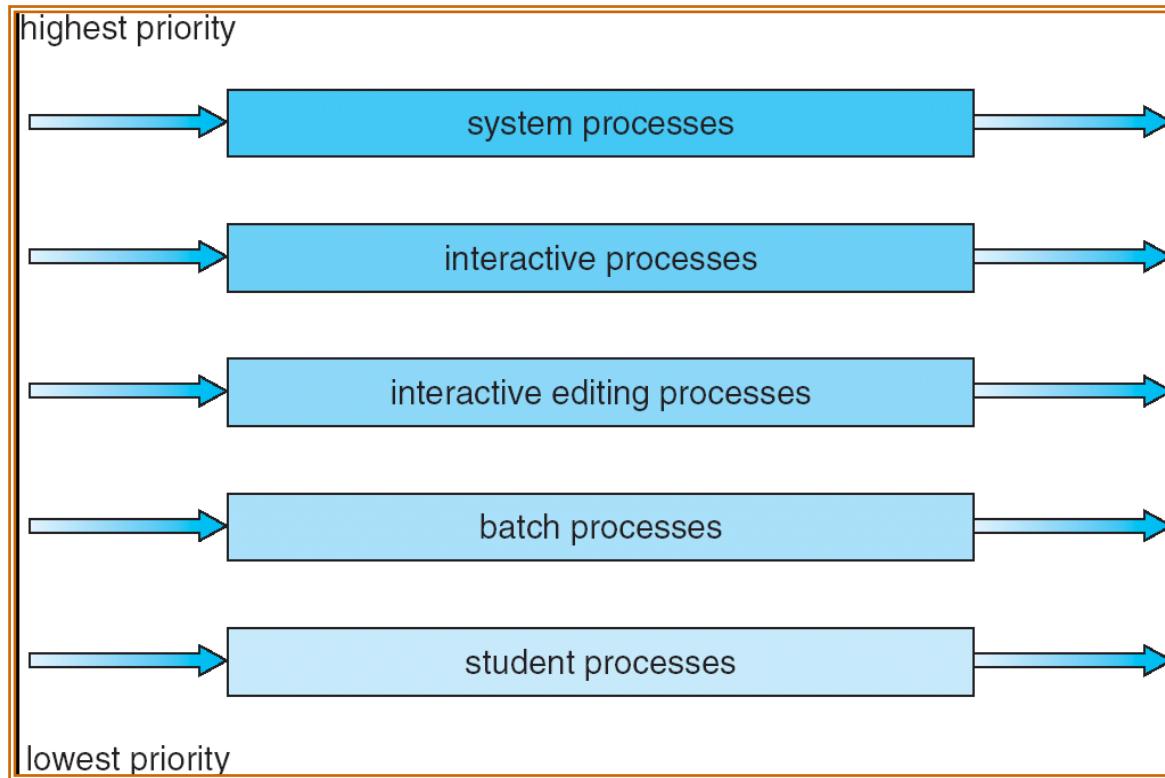
- foreground – RR (for low response time)
- background – FCFS (for simplicity and low context-switch overhead)

# Multilevel Queue

## How to schedule from among the queues?

- Fixed priority scheduling
  - i.e. the processes from foreground queue get the CPU, the background processes get the CPU only if the foreground queue is empty
  - Possibility of starvation.
- Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes
  - i.e., 80% to foreground queue, 20% to background queue
  - not necessarily optimal

# Multilevel Queue Scheduling



# Multilevel Feedback Queue

Idea:

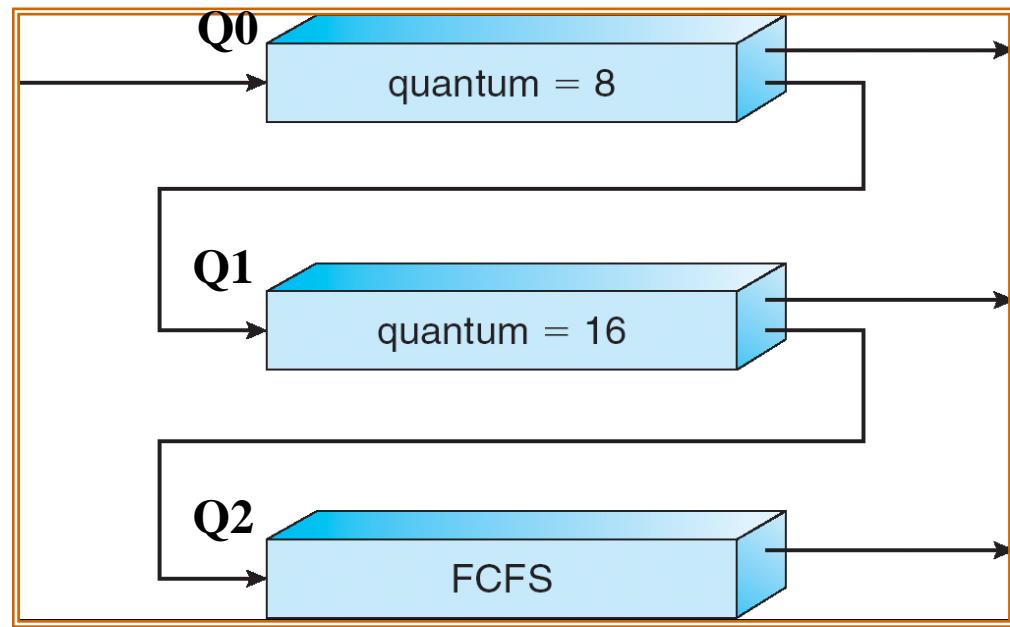
- **Use multilevel queues**
- **A process can move up and down queue hierarchy**
  - Why would a process move to a lower priority queue?
    - It is using too much CPU time
  - Why would a process move to a higher priority queue?
    - Has been starved of CPU for long time
    - A way to implement aging

# Multilevel Feedback Queue

- **Multilevel-feedback-queue scheduler defined by the following parameters:**
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service
- **This algorithm is the most general one**
  - It can be adapted to specific systems
  - But it is also the most complex algorithm

# Example of Multilevel Feedback Queue

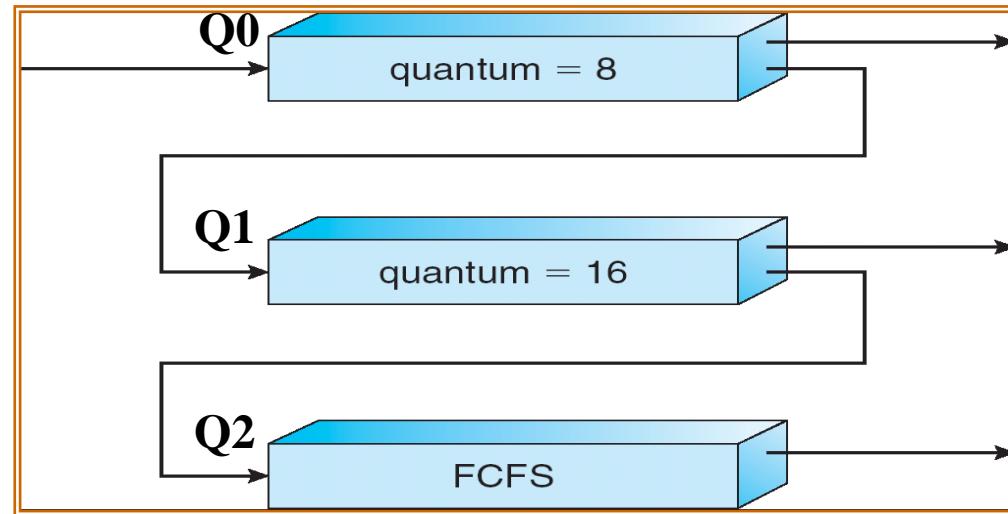
- Scheduler selects processes in Q0 first (highest priority)
  - If Q0 is empty, the processes from Q1 are selected.
  - If both Q0 and Q1 are empty, processes from Q2 are selected
- If a process arrives in a higher priority queue when another from a lower priority queue is running, the running process will be preempted, to allow the arriving process to run.
- When a process exhausts its quantum in either Q0 or Q1, it is preempted and moved to the lower priority queue.



# Example of Multilevel Feedback Queue

- **Scheduling example**

- A new job enters queue  $Q_0$  which is served FCFS. When it gains CPU, job receives 8 milliseconds.
- If it does not finish in 8 milliseconds, job is preempted and moved to queue  $Q_1$  and served again FCFS to receive another 16 additional milliseconds.
- If it still does not complete, it is preempted and moved to queue  $Q_2$ .



# Multilevel Feedback Queue Discussion

**Exact properties depend on the parameters**  
**Flexible enough to accommodate most requirements**

**The convoy example:**

- One process with long CPU burst time
- Several I/O bound processes with short CPU burst time
- Even if all processes start at the same level, the CPU-intensive process will be soon demoted to a low priority queue
- The I/O bound processes will remain at high priority and will be swiftly serviced, keeping the I/O devices busy

## In practice...

- The methods we have seen are all used in practice (except *pure SJF* which is impossible)
- Sophisticated OS provide the system manager with a library of methods, which it can choose and combine as needed after observing the behavior of the system.

# Also...

- Our study of scheduling methods is theoretical, does not consider in detail all the problems that arise in CPU scheduling.
- Eg. CPU schedulers cannot give the CPU to a process for all the time it needs
  - Because in practice, the CPU will often be interrupted by some external event before the end of its cycle.
- Also, this study does not consider the execution times of the scheduler at all.

# Summary of scheduling algorithms

- **First come, first served (FCFS)**
  - simple, low system time (overhead), low quality
- **Shortest Job First (SJF)**
  - Must know processing times (not practical)
  - Must predict using the exponential mean of the past
- **Scheduling with priority**
  - A class of algorithms
- **Round-Robin**
  - FCFS with preemption
- **Multilevel Queues**
  - Possible to use different algorithms with each queue
- **Multilevel Feedback Queues**
  - Combines several techniques

# Overview of advanced scheduling topics

- **Scheduling with several identical CPUs**
- **Evaluation model**

# Multiple-Processor Scheduling

Good news:

- **With multiple CPUs, we can share the load**
- **We can also share the OS overhead**

Bad news:

- **We are expected to efficiently share the load**
- **Managing the OS structures gets more complicated**
- **The scheduling gets more complex**

# Multiple-Processor Scheduling - Approaches

We assume *homogeneous processors*

- i.e. all processors have the same functionality

Still, the scheduler (and possibly the whole OS) might be run on a single CPU

- *Asymmetric multiprocessing* – only one CPU accesses the system data structures and makes scheduling decisions
  - alleviates the need for data sharing
  - but might become bottleneck
- *Symmetric multiprocessing (SMP)* – each CPU is self-scheduling
  - Either from a common ready queue, or from a private ready queue
  - Care must be taken when CPUs access common data structures
  - Virtually all modern operating systems support SMP including Windows XP, Solaris, Linux, Mac OS X.

# SMP Scheduling Issues

- **Processor Affinity**
  - When process runs on a physical CPU, cache memory is updated with content of the process
  - If the process is moved to another CPU, benefits of caching is lost.
  - SMP systems try to keep processes running on the same physical CPU – known as processor affinity
- **Load Balancing**
  - Required when each CPU has its own Ready Queue
  - Push migration: a task is run periodically to redistribute load among CPUs (their ready queue)
  - Pull migration: An idle CPU (i.e. with an empty ready queue), will pull processes from other CPUs
  - Linux supports both techniques
  - Can counteract the benefits of processor affinity, since it moves processes from one CPU to another

# Evaluation methods and comparison of algorithms

- **Deterministic modeling**
- **Queuing models**
- **Simulation**

# Deterministic modeling

- **Essentially, what we have already done by studying the behavior of several algorithms on several examples**

# Use of queuing theory

- **Analytical method based on probability theory**
- **Simplified model: in particular, the OS times are ignored**
- **However, it makes estimates possible**

# Queuing theory: Little's formula

- An important result:
- $n = \lambda \times W$ 
  - $n$ : average length of the queue
  - $\lambda$  : process arrival rate in queue
  - $W$ : average waiting time in the queue
- Eg.
  - $\lambda$  if processes arrive 3 per sec.
  - $W$  and they stay in the queue for 2 seconds
  - $n$  the average queue length will be ???
- Exercise: Solve also for  $\lambda$  and  $W$
- Observe that for  $n$  to be stable,  $\lambda \times W$  must be stable
  - If  $n$  must remain 6 and  $\lambda$  goes up to 4, what must be  $W$ ?

# Simulation

- **Build a (*simplified...*) model of the sequence of events in the OS**
- **Assign a duration of time to each event**
- **Assume a certain sequence of external events (e.g. arrival of process, etc.)**
- **Run the model for this sequence to get stats**

# Important points in this chapter

- **Ready Queue for CPU**
- **Scheduling criteria**
- **Scheduling algorithms**
  - FCFS: simple, not optimal
  - SJF: optimal, difficult implementation
    - exponential averaging
  - Priorities
  - Round-Robin: selection of the quantum
- **Algorithm evaluation, queuing theory,**
  - Little's formula

Thank You!

ຂໍ ດັບ ດຸນ

DMnvwd

Gracias

Dankie

Obrigado!

WAD MAHAD

SAN TAHAY

감사합니다.

Viel  
Dank



شکریا

Díky



Eυχαριστώ

Teşekkürler

Grazie

Bedankt

Köszönettel

謝謝

GADDA GUEY

Urakoze

Merci

مشکرم